



Scientific Days Cameroon  
*December 20, 2022*

# Efficient Buffer Overflow Mitigation In Virtualized Clouds Using Intel EPT-based Sub-Page Write Protection Support

Stella Bitchebe  
Ph.D. Candidate, Université Côte d'Azur & ENS Lyon, France  
Advisor: Pr. Alain Tchana  
*bitchebe@i3s.unice.fr*

# Context: Buffer Overflow in Virtualized Clouds

---

Why the Cloud?


- ▶ Attractive costs (e.g., *FaaS*)
- ▶ Management tasks simplification (e.g., *SaaS*, *IaaS*)
- ▶ Efficiency of Cloud computing

# Context: Buffer Overflow in Virtualized Clouds

Why the Cloud?


- ▶ Attractive costs (e.g., *FaaS*)
- ▶ Management tasks simplification (e.g., *SaaS*, *IaaS*)
- ▶ Efficiency of Cloud computing

Virtualized Cloud Infrastructure




# Context: Buffer Overflow in Virtualized Clouds

What is buffer overflow?



# Context: Buffer Overflow in Virtualized Clouds

What is buffer overflow?




Importance of Buffer Overflow

- 70% Google Chrome's bugs
- 70% of Microsoft vulnerabilities
- Top vulnerability in 2022

# State-of-the-art Techniques


---

- Canary




# State-of-the-art Techniques

## ● Canary




## ● Guard Page




# State-of-the-art Techniques

---

## ● Canary



## ● Guard Page



## ● Hardware Capabilities

CHERI: Capabilities Hardware Enhance RISC-V Instructions



# Secure Allocators: SlimGuard

---

## Overview

- State-of-the-art BIBOP (Big Bag Of Pages) secure allocator
- Published in 2019
- Has proven more memory-efficient than other state-of-the-art BIBOP allocators (Guarder, FreeGuard, etc.)
- Available and functional code

# Secure Allocators: SlimGuard

## Overview

- State-of-the-art BIBOP (Big Bag Of Pages) secure allocator
- Published in 2019
- Has proven more memory-efficient than other state-of-the-art BIBOP allocators (Guarder, FreeGuard, etc.)
- Available and functional code
- Fine-grained bag sizes (not size-of-two -2, 8, 16, ...- like others)



# Secure Allocators: SlimGuard

## Overview

- State-of-the-art BIBOP (Big Bag Of Pages) secure allocator
- Published in 2019
- Has proven more memory-efficient than other state-of-the-art BIBOP allocators (Guarder, FreeGuard, etc.)
- Available and functional code
- Fine-grained bag sizes (not size-of-two -2, 8, 16, ...- like others)




## Protection policy: applied at the bag scale

- Configurable proportion of guard pages  $P$
- $1/P$  allocation pages between 2 guard pages

# Secure Allocators: SlimGuard

## Limits


Significant memory overhead




# Secure Allocators: SlimGuard

## Limits

### Significant memory overhead





### Inefficient protection policy





# Secure Allocators: SlimGuard

## Custom Protection Policy




# Secure Allocators: SlimGuard

## Custom Protection Policy



# Synchronous Detection vs Memory Overhead

Canary: **modest memory overhead** + asynchronous overflow detection  
Guard page: significant memory consumption + **synchronous overflow detection**



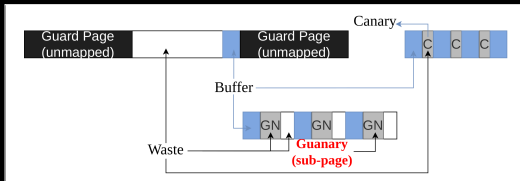


# Synchronous Detection vs Memory Overhead

Canary: modest memory overhead + asynchronous overflow detection


Guard page: significant memory consumption + synchronous overflow detection

GuaNary: **modest memory overhead + synchronous overflow detection**




# GuaNary: Canary and Guard Page Countermeasure

## Intel Sub-Page write Permission (SPP)




# GuaNary: Challenges

## $C_1$ : One size does not fit all



# GuaNary: Challenges


## $C_1$ : One size does not fit all



## $C_2$ : Costly hypercalls

# GuaNary: Challenges

**C<sub>1</sub>: One size does not fit all**




**C<sub>2</sub>: Costly hypercalls**

**C<sub>3</sub>: Page heterogeneity**

# GuaNary: Challenges


## $C_1$ : One size does not fit all



## $C_2$ : Costly hypercalls


## $C_3$ : Page heterogeneity

## $C_4$ : SPP Page heterogeneity



# GuaNary: Challenges



## $C_1$ : One size does not fit all



## $C_2$ : Costly hypercalls

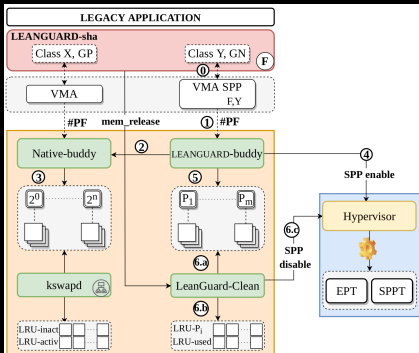
## $C_3$ : Page heterogeneity

## $C_4$ : SPP Page heterogeneity



# LeanGuard: GuaNary-based Secure Allocator

## Overview

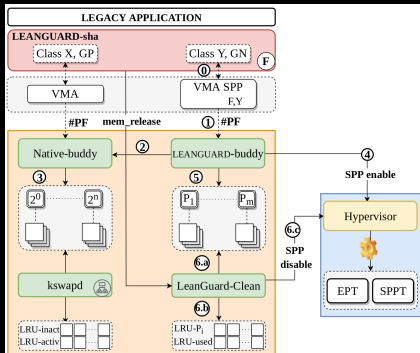


$c_1$ : use GuaNary or guard page according to the size of buffers



# LeanGuard: GuaNary-based Secure Allocator

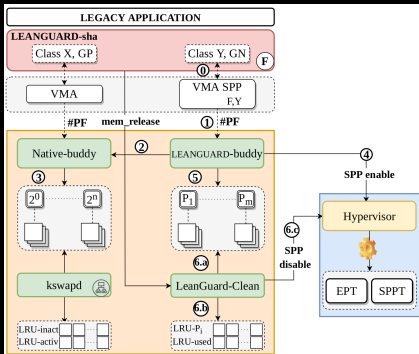
## Overview



- $c_1$ : use GuaNary or guard page according to the size of buffers
- $c_2$ : batching hypercalls for optimal performance
- $c_3$ : new buddy allocator (LeanGuard-buddy)

# LeanGuard: GuaNary-based Secure Allocator

## Overview



$c_1$ : use GuaNary or guard page according to the size of buffers

$c_2$ : batching hypercalls for optimal performance

$c_3$ : new buddy allocator (LeanGuard-buddy)

$c_4$ : formula to determine which pattern to use; 202 pools, one per pattern

# LeanGuard: GNaNary-based Secure Allocator

---

## Implementation

- ▶ Xen hypervisor version 4.10: 600LOCs addition

# LeanGuard: GNaNary-based Secure Allocator

---

## Implementation

- ▶ Xen hypervisor version 4.10: 600LOCs addition
- ▶ Linux OS version 5.11.14: 750LOCs addition

# LeanGuard: GNaNary-based Secure Allocator

---

## Implementation

- ▶ Xen hypervisor version 4.10: 600LOCs addition
- ▶ Linux OS version 5.11.14: 750LOCs addition
- ▶ SlimGuard secure allocator: 100LOCs addition

# Evaluations

---

## Benchmarks

- ▶ Micro-benchmark: 1GB working set application, random buffers allocation from all SlimGuard Classes
- ▶ Macro-benchmark: PARSEC suite

# Evaluations


## Memory Consumption Results



- ▶ GuaNary effectively reduces memory consumption compared to Canary and guard page
- ▶ To protect **50%** of the allocated buffers (F=2), LeanGuard on average uses **60%** less memory compared to SlimG+GP
- ▶ Using the same amount of memory as SlimG+GP, LeanGuard allows protecting **25×** more buffers than SlimG+GP

# Evaluations

## Performance Overhead




- ▶  $\sim 7.7\%$  overhead on average
- ▶ The main source of this overhead is the SPP page table walk on TLB miss



# Evaluations

---

## Scalability



- LeanGuard-buddy uses a lock mechanism

# Conclusion

---

- ▶ We introduce a novel technique to mitigate buffer overflow in virtualized clouds: **GuaNary**
  - ▶ based on Intel sub-page write permission (SPP)
  - ▶ **modest memory overhead + synchronous overflow detection**
- ▶ We propose
  - ▶ LeanGuard: a complete software stack that uses GuaNary
  - ▶ LeanGuard-secure allocator:
    - consumes **8.3×** less memory compared to the SlimGuard allocator (up-to-date state-of-the-art)
    - Allows protecting **25×** more buffers than SlimGuard for the same amount of memory